



# Docker From The Ground Up

## Part 1

A whistle-stop tour of containerisation, Docker images, Dockerfiles, sharing images, running containers and relevant aspects of best practice.



# Who's This Guy?

**Matt Todd**

Principal Architect @ Resonate.tech

Cloud Native Advocate

`github.com/mattjtodd`

`hub.docker.com/u/mattjtodd`





**Thanks to our Sponsor!**

**BlackCat /**





**DID SOMEONE SAY**



**BEER?**





The dawn of containers..... (Waaaaayyyyyy Pre DevOps)

---

**Chroot**



# Chroot (Jail)

- Restricts a processes (and all children's) view of a servers FS to a specific portion
- Maps a process root FS to some other directory
- Prevents access to outside of the new root
- Introduced in **1979** Unix V7 the BSD **1982**





---

**Fast Forward to 2000.....**



# FreeBSD Jails

- Allows administrators to partition a [FreeBSD](#) computer system into several independent, smaller systems – called “jails”
- Ability to assign an IP address for each system and configuration.



## 2001 - 2005

- 2001 - Linux VServer
- 2004 - Solaris COnainers
- 2005 - Open Virtuozzo



## 2006 - Process Containers

- Launched by Google in 2006) was
- designed for limiting, accounting and isolating resource usage (CPU, memory, disk I/O, network) of a collection of processes.
- It was renamed “Control Groups (cgroups)” a year later
- Merged to Linux kernel 2.6.24.



## 2008 LXC (Linux Containers)

- First, most complete implementation of Linux container manager
- Implemented in 2008 using cgroups and Linux namespaces
- Works on a single Linux kernel without requiring any patches
- First “Standard” container platform



# Warden & LMCTFY

- CloudFoundry started Warden in 2011
- using LXC in the early stages and later replacing it with its own
- Warden can isolate environments on any operating system, running as a daemon
- providing an API for container management.
- It developed a client-server model to manage a collection of containers across multiple hosts, and Warden includes a service to manage cgroups, namespaces and the process life cycle.
- [Let Me Contain That For You](#) (LMCTFY) kicked off in 2013 as an open-source version of Google's container stack, providing Linux application containers.
- Applications can be made “container aware,” creating and managing their own subcontainers.
- Active deployment in LMCTFY stopped in 2015 after Google started contributing core LMCTFY concepts to libcontainer, which is now part of the [Open Container Foundation](#).





# Container Process Isolation

## Namespaces

- Resource allocation / isolation
  - Mount (Filesystem)
  - Process (PID)
  - User ID
  - Networks (Virtual Devices)
  - ...

## CGroups

- Resource *usage* limits & Prioritisation
  - Memory
  - CPU DiskIO
  - ...



Solomon Hykes




# Docker container engine

- LXC
- Libcontainer (<https://github.com/docker/libcontainer>)
- RunC (<https://github.com/opencontainers/runc>)
- Donated to Open Container Initiative 2015
  - <https://www.opencontainers.org/>
  - <https://github.com/opencontainers/runc>

---

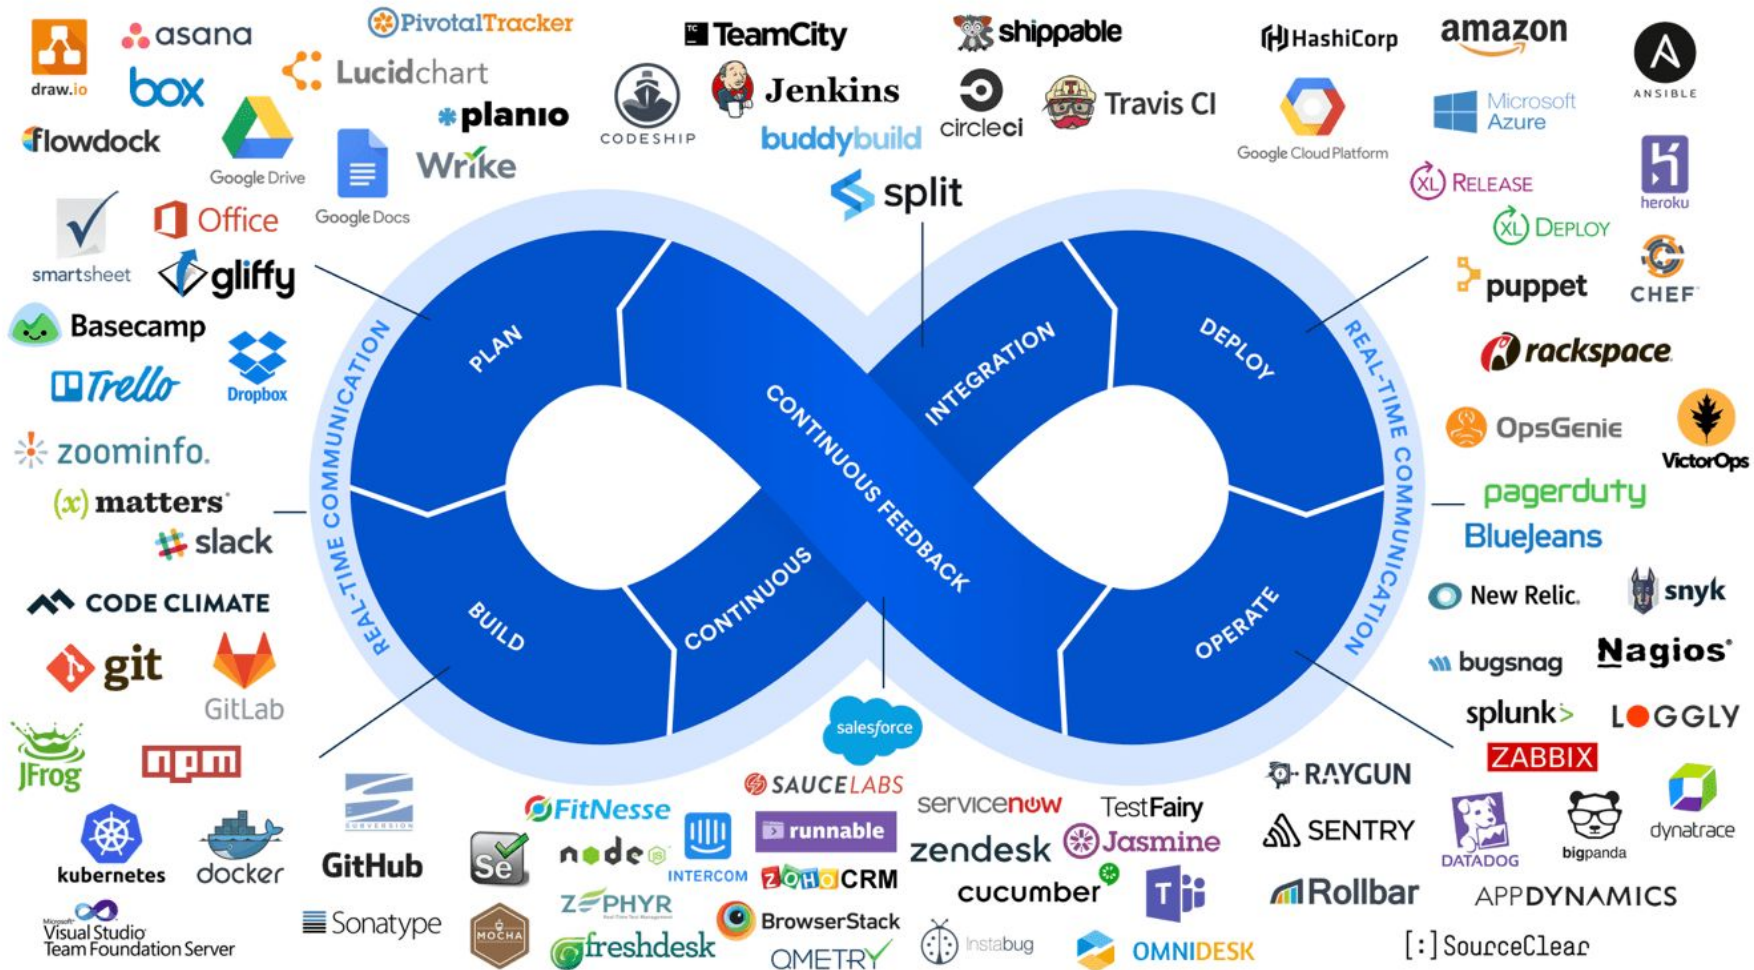
# Build, Ship, Run






# Docker / Containerisation

- Runtime Standardisation
- Scalability
- Dependency Management & Versioning
- Lightweight & Resource Isolated
- Shared FS Layering
- Repeatability of Builds
- Portability
- Cloud Native (CNCF)
- Immutability (Cattle Not Pets)
- Rich Sharing Ecosystem



# DevOps and Containers

Creation, deployment, and management

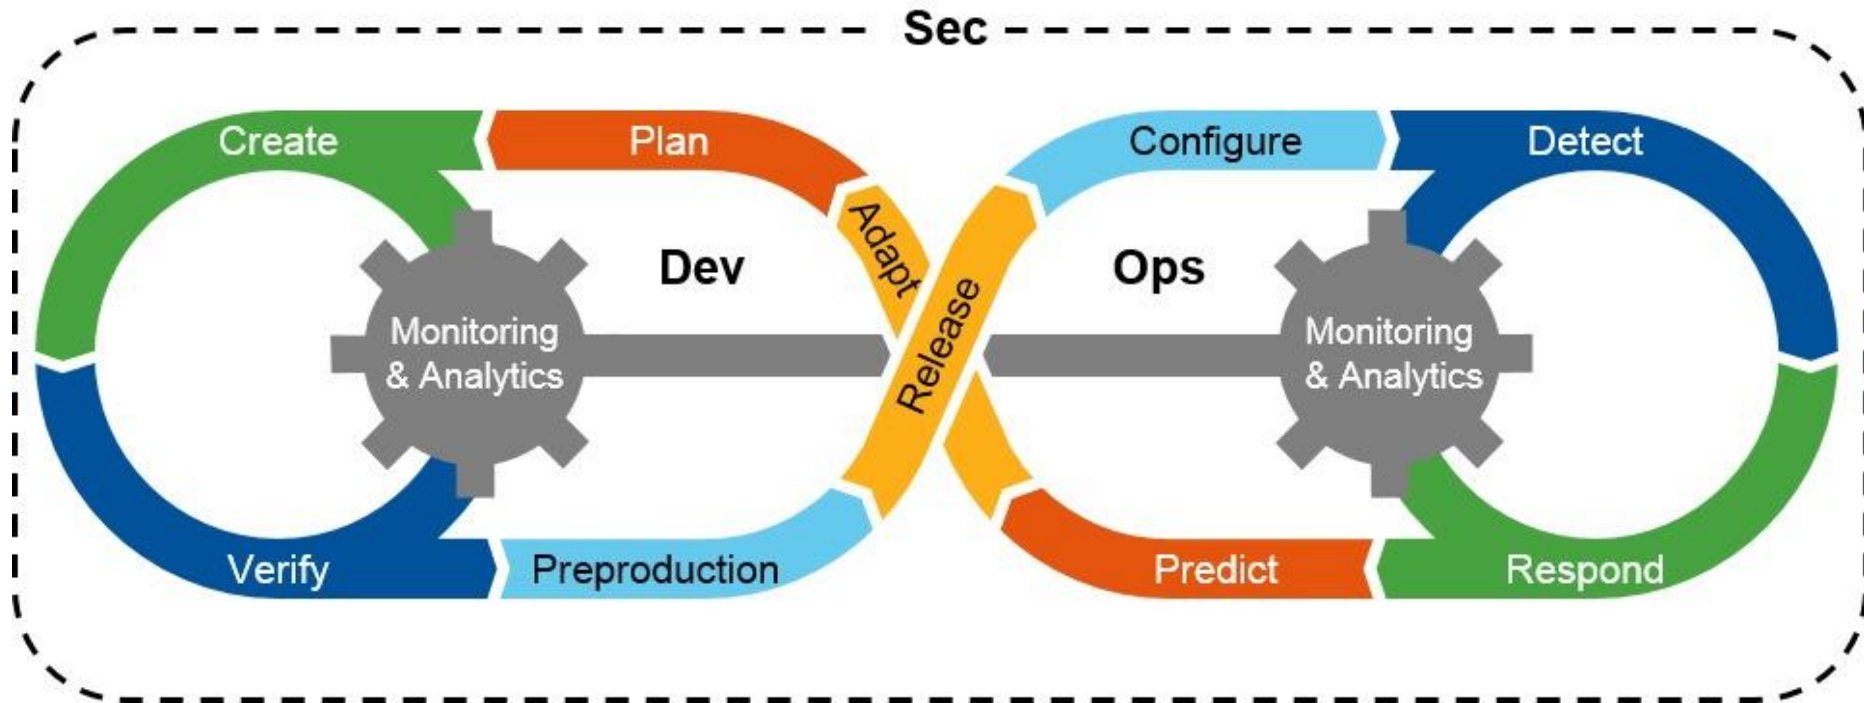


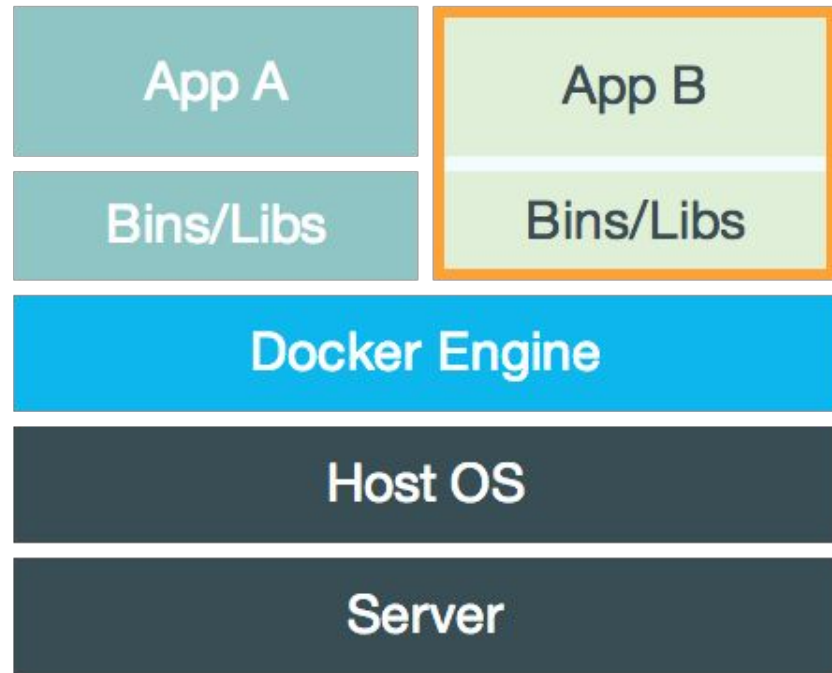
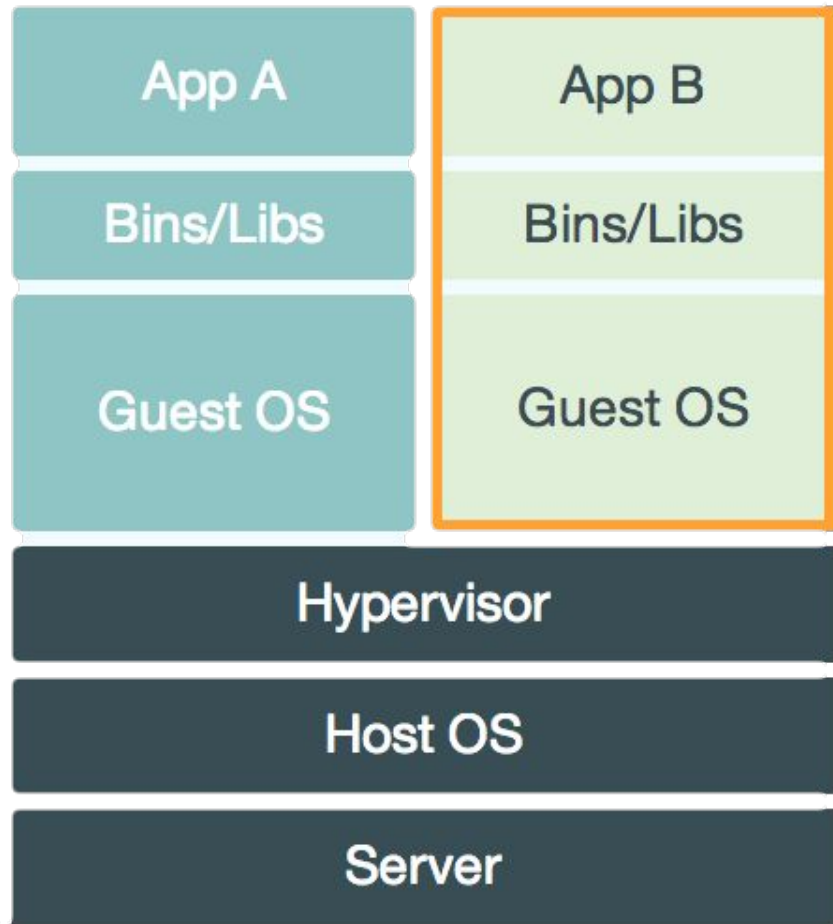


---

# Shift Left For Security

Sec







**WORKED FINE IN  
DEV**

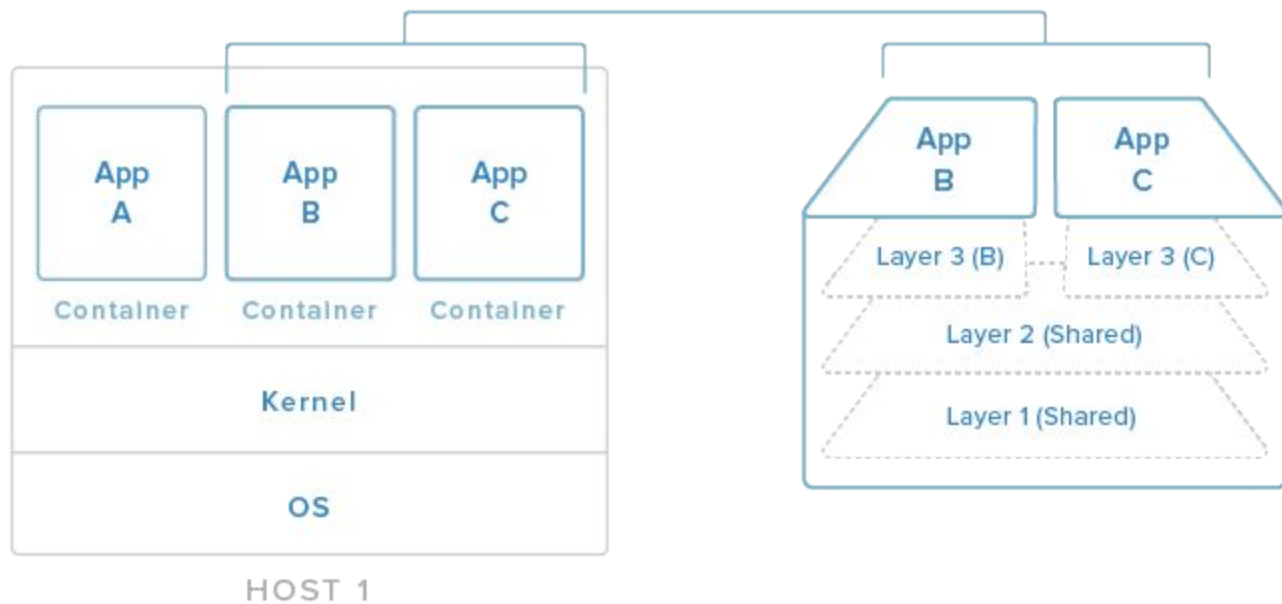
**OPS PROBLEM NOW**



# Container Use Cases


- Applications
- Databases
- Services
- Build Agents
- Test harnesses
- POCs
- X11 services (Jessie Frazelle) <https://blog.jessfraz.com/>)
- .... It's just a process after all....

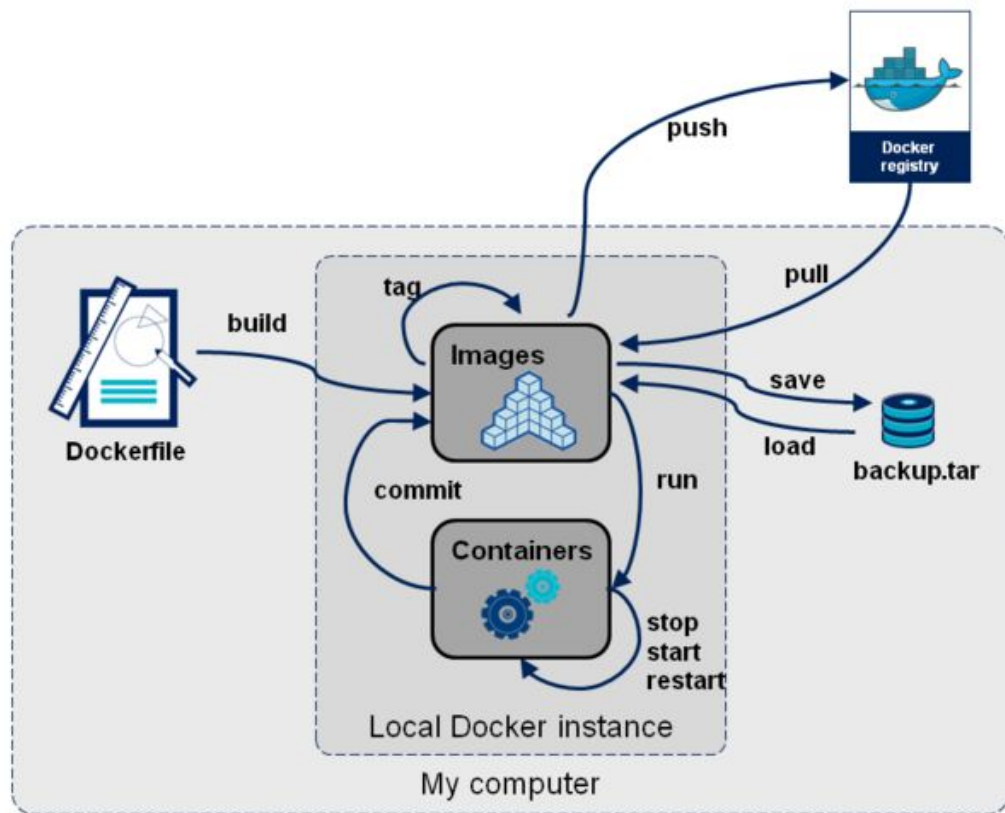
## CONTAINER OVERVIEW





# DOCKER COMPONENTS

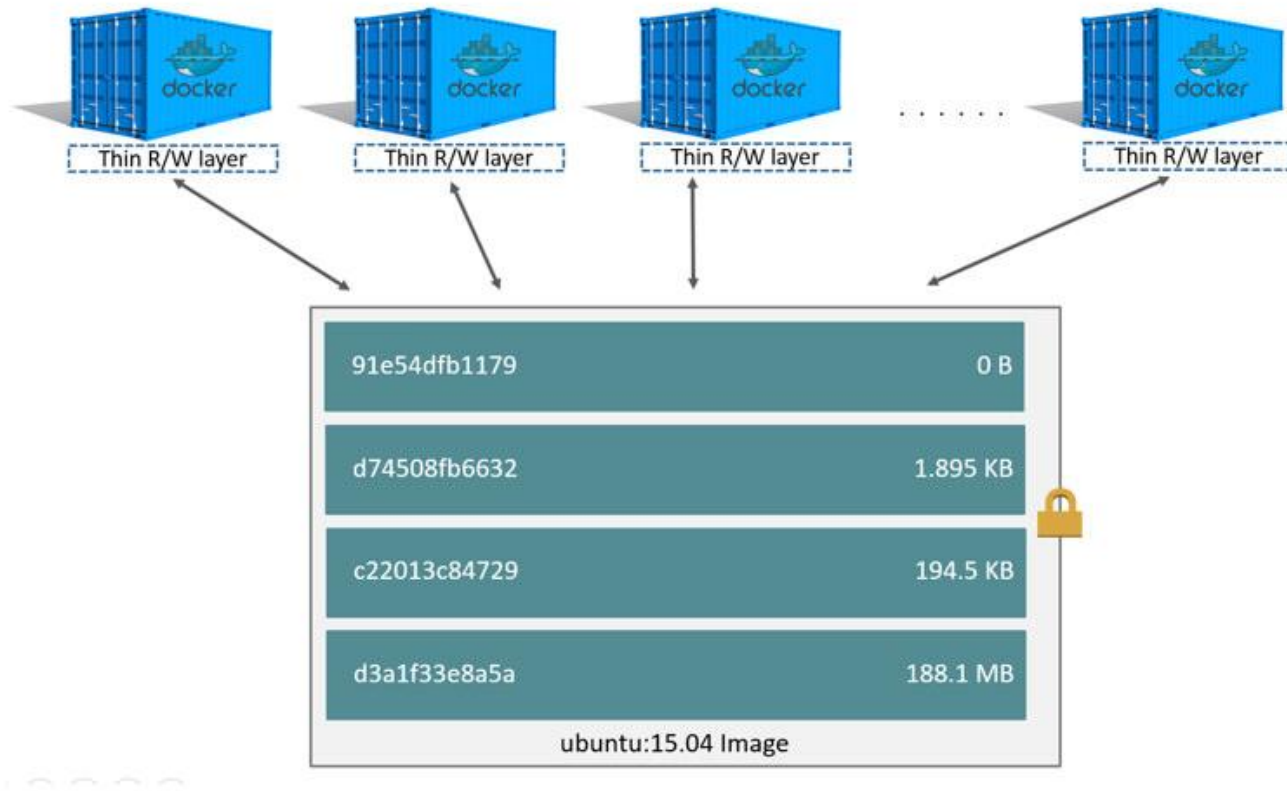






# Docker Daemon CLI Communication

- Non networked unix file socket (*/var/run/docker.sock*)
- HTTP / HTTPS (Please configure TLS!)
- SSH tunnel port forward socket.
- *-H* param can be used to configure this



Union Filesystem

# Union File System



All layers except the topmost are readonly



# Dockerfiles

- Builds layers in the union FS
- SHA-256 digest of the layer result of the directive
- Build cache for repeat build efficiency



## Simple Example

```
FROM alpine:3.8
```

```
COPY greeting.txt ./
```

```
USER nobody
```

```
CMD [ "cat", "greeting.txt" ]
```





# Simple Example

```
docker build -t hello-docker-birmingham .
```

```
Sending build context to Docker daemon  3.072kB
```

```
Step 1/4 : FROM alpine:3.8
```

```
---> 11cd0b38bc3c
```

```
Step 2/4 : COPY greeting.txt ./
```

```
---> 6e621934431d
```

```
Step 3/4 : USER nobody
```

```
---> Running in 36343a423445
```

```
Removing intermediate container 36343a423445
```

```
---> 0d89fbf00a6c
```

```
Step 4/4 : CMD ["cat", "greeting.txt"]
```

```
---> Running in 97c4c7d01da4
```

```
Removing intermediate container 97c4c7d01da4
```

```
---> 6913f6e1fd4b
```

```
Successfully built 6913f6e1fd4b
```



# What's going on?

- The base image is Alpine (a tiny linux distro)
- Copy a file from the build context dir
- As the user nobody (step-down from root)
- Run the cat command with the file



# Docker Metadata

- Most constructs have the ability to be *inspected*.
- For example:
  - `docker image inspect hello-docker-birmingham`
- Use *jq* for piping transforms
- Alternatively go formatters which can be CLI provided or in
  - `~/ .docker/config.json`



## Fire it Up!

```
$ docker run hello-docker-birmingham
```

```
Hello Docker Birmingham!%
```



# How about a webserver?

```
FROM alpine:3.8
```

```
RUN apk add --no-cache npm
```

```
RUN npm install http-server -g
```

```
WORKDIR /etc/files
```

```
COPY greeting.txt ./
```

```
USER nobody
```

```
CMD ["http-server", "./"]
```



# How about a webserver?

```
$ docker build -t http-server .
```

```
$ docker history http-server
```

IMAGE SIZE	CREATED COMMENT	CREATED BY
a2ad731ff8ba 0B	4 hours ago	/bin/sh -c #(nop) CMD ["http-server" "./"]
ceacd9e50527 0B	4 hours ago	/bin/sh -c #(nop) USER nobody
4f707fe61903 12B	4 hours ago	/bin/sh -c chown nobody:nobody *
af8d819dae00 12B	4 hours ago	/bin/sh -c #(nop) COPY file:2c0ebb879350f3ff...
797657cae401	4 hours ago	/bin/sh -c #(nop) WORKDIR /etc/files



# How about a webserver?

```
$ docker run -name server -p 127.0.0.1:9090:8080 http-server
```

```
Starting up http-server, serving ./
```

```
Available on:
```

```
    http://127.0.0.1:8080
```

```
    http://172.17.0.2:8080
```

```
Hit CTRL-C to stop the server
```

```
$ curl localhost:9090/greeting.txt
```

```
Hello World!%
```



# What's going on?

- Webserver process is being started
- The webserver's port 8080 is being made available to localhost *only*
- When the service terminates the container definition is deleted
- Let's leave it running for now.....







# Docker Containers

- Debugging
- Instrumentation
- Resource Constraints
- Labels
- Volumes
- Process Managers
- Logs



# Debugging

- `docker exec -it server top`
- You can invoke any app installed in the container FS and you have permission to do so:
  - Install Packages using your ***favourite*** package manager
  - List process `ps -ax`
  - ....
- Will work on the ***writable*** layer of the container\*
- No need for ***sshd!***



# Container Stats

- CLI Stats / Endpoint
- Prometheus Endpoint (<https://docs.docker.com/config/thirdparty/prometheus/>)
- Health

```
$ docker stats
```



# Labels

- Can be added to both images and containers
- Consider using the Label Schema (<https://label-schema.org/rc1/>)
- Immensely useful when running in production!



# Resource Constraints

- CPU
- Memory
- DiskIO
- BlkIO

```
docker run --rm -it --cpus=0.1 -m 80m tomcat:9
```

```
docker run --rm -it --cpus=1 -m 10m tomcat:9
```



*I do not introduce the log.*

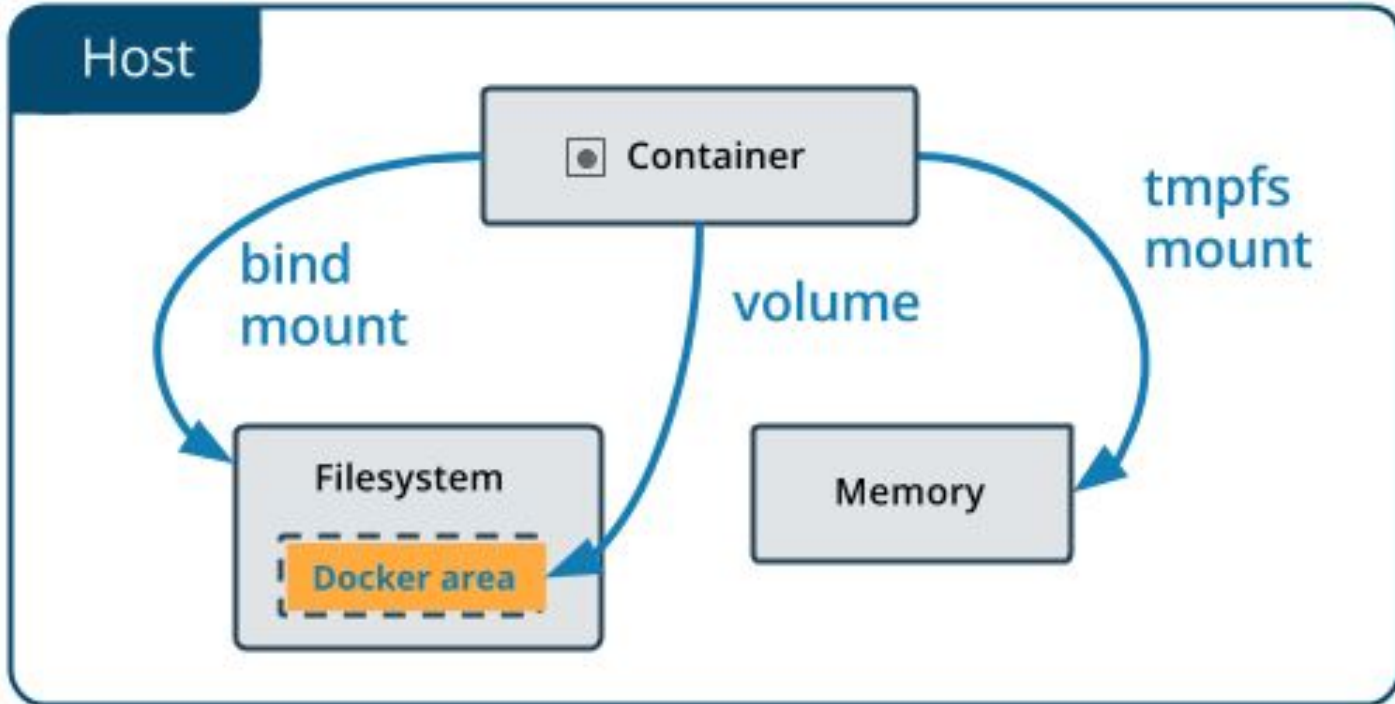


# Logs

- Stdout / Stderr are logged to the configured driver
- Configurable logging drivers
  - Json-file (default)
  - FluentD
  - Gelf
  - Splunk
  - Awslogs (Cloud Watch)
  - ...
- For json-file you can use
  - `docker logs -f server`
- For other drivers, check the `agg. service!`



Vo





# Volumes

- Physical (Bind)
- Logical (Named)
- Drivers available
  - Local (only visible on host)
  - NFS
  - Azure
  - Portworx
  - REX-Ray
  - ... Lots more open source!



# Bind Volumes

- Mounts are visible to the writable container layer to the filesystem

```
$ docker run --rm -v $PWD:/etc/files:ro -p 9091:8080 http-server
```

- Watch your permissions when using **:rw**!
- Portability issues between hosts
- Great for build agents!



# Named Volumes

- Are managed and created and linked to zero or more containers
- Like virtual USB sticks!
- Persist beyond container restarts
- Managed via pruning

```
$ docker run --rm -v http-server:/etc/files:ro -p 9091:8080
```

```
http-server
```



Docker Hub & Publishing



# Publishing

- Docker Hub!
- “Official” Repos
- Private Repos, e.g. Nexus, Artifactory
- Image Signing
- Image scanning
- Images must be *tagged* with the repo coordinates

```
$ docker build -t http-server -t mattjtodd/http-server .
```

```
$ docker tag http-server mattjtodd/http-server
```

```
$ docker login && docker push mattjtodd/http-server
```



Whoaaaaa!



# Some Best Practices

- Label all the things! (<https://label-schema.org/rc1/>)
- Take care to understand your process owner
- User a process manager (tini, gosu and --init)
- Don't bloat; alpine where possible avoid unnecessary layers esp. Subsequent deletes
- Use the image directive cache *wisely*...
- Be careful when mapping to
  - Bind volume mounts
  - Ports to `0.0.0.0:<port>`
- Whitelist the build context with `.dockerignore`



<https://www.katacoda.com/courses/docker/>

---



Get ready for next time.....



# Homework!

- Install Docker
- Pull images from the repo for next time
- I'll post which are needed before the next session
- Ask on the Slack channel if you have any issues / questions